

importing dependencies

```
In [2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, precision_score
```

Data collection and Processing

```
In [17]: titanic_data=pd.read_csv("train.csv")
titanic_data.head()
```

Out[17]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500

```
In [18]: titanic_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
 #   Column        Non-Null Count  Dtype
---  -
 0   PassengerId   891 non-null    int64
 1   Survived      891 non-null    int64
 2   Pclass        891 non-null    int64
 3   Name          891 non-null    object
 4   Sex           891 non-null    object
 5   Age           714 non-null    float64
 6   SibSp         891 non-null    int64
 7   Parch         891 non-null    int64
 8   Ticket        891 non-null    object
 9   Fare          891 non-null    float64
10   Cabin         204 non-null    object
11   Embarked      889 non-null    object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
```

```
In [19]: titanic_data.shape
```

```
Out[19]: (891, 12)
```

```
In [20]: titanic_data["Pclass"].value_counts()
```

```
Out[20]: Pclass
3      491
1      216
2      184
Name: count, dtype: int64
```

```
In [21]: titanic_data.isnull().sum()
```

```
Out[21]: PassengerId    0
Survived              0
Pclass                0
Name                  0
Sex                   0
Age                  177
SibSp                 0
Parch                 0
Ticket                0
Fare                  0
Cabin                 687
Embarked              2
dtype: int64
```

Handling the missing values

```
In [22]: titanic_data=titanic_data.drop(columns="Cabin",axis=1)
```

```
In [26]: titanic_data.describe()
```

Out[26]:

	PassengerId	Survived	Pclass	Age	SibSp	Parch	Fare
count	891.000000	891.000000	891.000000	714.000000	891.000000	891.000000	891.000000
mean	446.000000	0.383838	2.308642	29.699118	0.523008	0.381594	32.204208
std	257.353842	0.486592	0.836071	14.526497	1.102743	0.806057	49.693429
min	1.000000	0.000000	1.000000	0.420000	0.000000	0.000000	0.000000
25%	223.500000	0.000000	2.000000	20.125000	0.000000	0.000000	7.910400
50%	446.000000	0.000000	3.000000	28.000000	0.000000	0.000000	14.454200
75%	668.500000	1.000000	3.000000	38.000000	1.000000	0.000000	31.000000
max	891.000000	1.000000	3.000000	80.000000	8.000000	6.000000	512.329200

In [27]:

titanic_data['Age'].fillna(titanic_data['Age'].mean(),inplace=True)

In [29]:

titanic_data['Embarked'].mode()

Out[29]:

0 S
Name: Embarked, dtype: object

In [38]:

titanic_data['Embarked'].fillna(titanic_data['Embarked'].mode()[0],inplace=True)

In [39]:

titanic_data.isnull().sum()

Out[39]:

PassengerId 0
Survived 0
Pclass 0
Name 0
Sex 0
Age 0
SibSp 0
Parch 0
Ticket 0
Fare 0
Embarked 0
dtype: int64

Data Analysis

In [43]:

titanic_data['Survived'].value_counts()

Out[43]:


Survived
0 549
1 342
Name: count, dtype: int64

In [44]:

sns.set()


```
In [48]: sns.countplot(x="Survived",data=titanic_data,hue="Survived")
```

```
Out[48]: <Axes: xlabel='Survived', ylabel='count'>
```




```
In [53]: sns.countplot(x="Sex",data=titanic_data,hue="Survived")
```

```
Out[53]: <Axes: xlabel='Sex', ylabel='count'>
```




```
In [54]: sns.countplot(x="Pclass", data=titanic_data, hue="Survived")
```

```
Out[54]: <Axes: xlabel='Pclass', ylabel='count'>
```



```
In [55]: sns.countplot(x="Embarked",data=titanic_data,hue="Survived")
```

```
Out[55]: <Axes: xlabel='Embarked', ylabel='count'>
```



One_hot Encoding

```
In [56]: titanic_data['Sex'].value_counts()
```

```
Out[56]: Sex
male      577
female    314
Name: count, dtype: int64
```

```
In [57]: titanic_data['Embarked'].value_counts()
```

```
Out[57]: Embarked
S      646
C      168
Q       77
Name: count, dtype: int64
```

```
In [61]: titanic_data.replace({'Sex':{'male':0, "female":1}, "Embarked":{"S":0, "C":1, "Q":2}},i
```

```
In [62]: titanic_data.head(3)
```

Out[62]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Embarked
0	1	0	3	Braund, Mr. Owen Harris	0	22.0	1	0	A/5 21171	7.2500	S
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	1	38.0	1	0	PC 17599	71.2833	C
2	3	1	3	Heikkinen, Miss. Laina	1	26.0	0	0	STON/O2. 3101282	7.9250	S

In [63]:

```
X= titanic_data.drop(columns=["PassengerId", "Name", "Ticket", "Survived"],axis=1)
Y= titanic_data["Survived"]
```

In [65]:

```
X
```

Out[65]:

	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked
0	3	0	22.000000	1	0	7.2500	0
1	1	1	38.000000	1	0	71.2833	1
2	3	1	26.000000	0	0	7.9250	0
3	1	1	35.000000	1	0	53.1000	0
4	3	0	35.000000	0	0	8.0500	0
...
886	2	0	27.000000	0	0	13.0000	0
887	1	1	19.000000	0	0	30.0000	0
888	3	1	29.699118	1	2	23.4500	0
889	1	0	26.000000	0	0	30.0000	1
890	3	0	32.000000	0	0	7.7500	2

891 rows × 7 columns

In [67]:

```
Y
```

```
Out[67]: 0      0
         1      1
         2      1
         3      1
         4      0
         ..
        886     0
        887     1
        888     0
        889     1
        890     0
        Name: Survived, Length: 891, dtype: int64
```

```
In [69]: x_train,x_test,y_train,y_test=train_test_split(X,Y,test_size=0.2, random_state=2)
```

```
In [71]: print(X.shape,x_train.shape,x_test.shape)

(891, 7) (712, 7) (179, 7)
```

Model training

```
In [72]: model=LogisticRegression()
```

```
In [73]: model.fit(x_train,y_train)
```

c:\Users\Soumaya\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\linear_model_logistic.py:469: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

n_iter_i = _check_optimize_result(

```
Out[73]: ▾ LogisticRegression ⓘ ?
          LogisticRegression()
```

Model Evaluation

```
In [74]: x_train_predictions=model.predict(x_train)
```

```
In [75]: print(x_train_predictions)
```



```
[0 1 0 0 0 0 0 1 0 0 0 1 0 0 1 0 1 0 0 0 0 0 1 0 0 1 0 0 1 0 1 1 0 0 1 0 1
0 0 0 0 0 0 0 1 1 0 0 1 0 1 0 1 0 0 0 0 0 0 0 1 0 1 0 0 1 1 0 0 1 1 0 1 0 0 1
0 0 0 0 0 0 0 1 0 0 0 1 0 0 0 1 0 1 0 0 1 0 0 0 1 1 1 0 1 0 0 0 0 0 1 0 0 0
1 1 0 0 1 0 0 1 0 0 1 0 0 1 0 1 0 1 0 1 0 1 1 1 1 1 1 0 0 1 1 1 0 0 1 0 0
0 0 0 0 1 0 1 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 1 1 0 0 1 0 1 0 1 1 1
0 0 0 1 0 0 0 1 0 0 1 0 0 0 1 1 0 1 0 0 0 0 0 0 1 1 0 1 1 1 1 0 0 0 0 0 0 0
0 1 0 0 1 1 1 0 0 1 0 1 1 1 0 0 1 0 0 0 0 1 0 0 0 1 0 0 0 1 0 1 0 1 0 0 0
0 0 0 0 0 0 1 0 1 0 0 1 0 0 1 0 1 0 1 1 0 0 0 0 1 0 1 0 0 1 0 0 0 1 0 0 0
0 1 1 0 0 0 0 0 0 1 0 1 0 0 0 0 0 1 1 1 0 0 0 1 0 1 0 0 0 0 0 0 1 1 0 1 1
0 1 1 1 0 0 0 0 0 0 0 0 0 0 1 0 0 1 1 1 0 1 0 0 0 0 1 1 0 0 0 1 0 1 1 1 0 0
0 0 1 0 0 0 1 1 0 0 1 0 0 0 0 1 0 0 0 0 0 1 0 0 0 0 1 0 1 1 1 0 1 1 0 0 0
0 1 0 1 0 0 1 1 0 0 0 0 1 0 0 0 0 1 1 0 1 0 1 0 0 0 0 0 1 0 0 0 0 1 1 0 0
1 0 1 0 0 1 0 0 0 0 0 0 0 0 0 1 0 0 1 1 0 0 0 1 1 0 1 0 0 1 0 0 0 1 1 0 1 0
0 0 0 0 1 0 0 1 0 1 1 0 0 1 0 0 1 0 0 0 1 0 1 1 0 0 1 1 0 1 0 1 1 1 0 1 0
0 1 0 0 1 0 0 1 0 0 0 0 1 1 0 0 1 0 1 0 0 0 0 0 0 1 1 1 0 0 1 1 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 1 1 0 0 0 0 0 0 0 0 0 0 0 0 1 0 1 0 0 0
0 0 1 0 0 0 0 0 1 0 1 0 1 0 0 0 1 0 1 1 1 0 0 0 1 0 1 0 0 0 1 1 1 0 0 1 1
0 0 0 1 0 1 0 0 0 0 0 1 1 0 1 1 1 0 0 0 1 0 0 0 0 1 0 0 0 1 0 0 1 0 0 0 0
1 0 0 1 0 1 0 0 0 1 1 1 1 1 0 0 1 1 0 1 1 1 1 0 0 0 1 1 0 0 1 0 0 0 0 0
0 0 0 1 1 0 0 1 0]
```

```
In [78]: training_data_accuracy=accuracy_score(y_train,x_train_predictions)
print(training_data_accuracy)
```

0.8075842696629213

```
In [79]: x_test_predictions=model.predict(x_test)
```

```
In [81]: print(accuracy_score(y_test,x_test_predictions))
```

0.7821229050279329

```
In [82]: model.score(x_test,y_test)*100
```

Out[82]: 78.2122905027933

```
In [83]: x_train.head(3)
```

Out[83]:

	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked
30	1	0	40.0	0	0	27.7208	1
10	3	1	4.0	1	1	16.7000	0
873	3	0	47.0	0	0	9.0000	0